

Agentic Synthesis against Counterexample-Supplemented Sketches

Muness Castle

Eric Rubeck

July 1, 2026

Abstract

In the RosterSynth worked example, the bad repair is easy to write: append -40 hours and the roster balances. The gate has to reject it anyway. The correct repair cancels duplicate booking 1802, because the operation matters, not only the final hour total. The method puts that distinction in the repository—sketch, corpus, code, prompt, and gate—instead of leaving it in a prompt, chat transcript, or subject-matter expert’s memory.

We use the term *counterexample-supplemented sketch* for the repository artifact and *agentic synthesis* for the loop that edits against it. A human writes a partial program-like sketch. A finite corpus stores observations and promoted counterexamples. Subject-matter experts correct concrete outputs. A coding agent edits deterministic code and prompt surfaces under known-code anchors. A gate blocks completion until replay and semantic comparison pass for every promoted case.

The loop borrows from sketching and counterexample-guided inductive synthesis (CEGIS), but the execution surface is different. Classical sketching and CEGIS can make formal claims when the search space, specification, and checker are formal. Here, the synthesizer is an IDE-shaped coding agent whose edits span ordinary source code and LLM prompt surfaces. The claim stays bounded: when the gate passes, the current strategy satisfies the repository’s encoded replay and comparison predicates for the promoted corpus.

The originating deployment ran inside a proprietary enterprise harness: a custom Cursor extension, `cursor://`-style commands, and an embedded web application where subject-matter experts loaded production examples, corrected outputs, and helped turn those corrections into input/output specifications. The public companion repository provides the sanitized RosterSynth slice: the sketch, fixtures, code, tests, and generated evidence needed to inspect the disclosed loop.

1 Introduction


Start with booking 1802. The roster has twin active 40-hour bookings on the pay-window end date. A naive repair can append -40 hours and make the hours math come out. Replay passes because the hours gap closed. The repair is still wrong: duplicate policy says cancel the higher duplicate booking. Compare has to check the booking ID and operation, not only the arithmetic.

The same shape shows up outside rosters. A parser can pass the easy cases while treating every colon as syntax. A data-cleaning rule can pass a local fixture while erasing a client-specific exception. The agent followed the signal it was given and missed the rule the repository had not yet made durable.

Tests still matter. They catch shape errors, exceptions, bad deltas, and regressions. The missing piece is the rule behind the promoted failure: why this tempting repair was wrong, which sketch clause changed, and what a future agent must preserve. Run-local prompts have the opposite weakness. They can carry intent quickly, but they fork, drift, and escape the repository’s checks.

The repository needs an artifact that records both the human strategy and the examples that sharpen it. The ingredients are:

1. a **sketch** $\mathcal{S}$: a partial program in prose and code-shaped structure;
2. a **counterexample-supplemented corpus** $\mathcal{E}$: examples, observations, and promoted failures;
3. **known-code anchors** $\mathcal{K}$: codebase-local shapes, APIs, conventions, and reference paths the agent must preserve;
4. a **dual-oracle implementation surface**: deterministic code for encodable holes and prompt-mediated completion for narrative holes;
5. a **gate** $\mathcal{G}$: replay plus semantic comparison over every promoted case.



The loop turns a tempting wrong repair into durable repository guidance.

The sketch records the human’s current strategy: which repairs are permitted, prioritized, or forbidden. The corpus records cases that have earned specification authority through promotion. The gate is the executable boundary for the current done state. A failure becomes valuable when it changes the sketch and the corpus so future agent passes inherit the rule instead of rediscovering the same local trap.

Central claim. A workflow is inspectable when every promoted counterexample connects four things: the sketch clause it supplements, the implementation or prompt surface that handles it, the replay check that verifies state repair, and the semantic compare check that rejects the tempting wrong repair. That gives maintainers a bounded gate invariant for the current promoted corpus and a promotion procedure for new failures.

Positioning. The comparison target is bounded. Existing agentic workflows can make tests and prompts durable. Classical CEGIS already has counterexample discipline when its formal assumptions hold. The repository setting here is smaller and messier: the human strategy, the promoted counterexample, and the executable check often live in separate artifacts. The work is to join them without pretending the result is a solver-backed proof.

Contributions. The contribution is a formulation plus a worked artifact.

1. a repository-native artifact: the counterexample-supplemented sketch;
2. a process model that turns subject-matter expert corrections into input/output specs, counterexamples, and golden data;
3. two hole-filling roles: Oracle A, deterministic code the agent maintains, and Oracle B, prompt-mediated completion for sketch-declared narrative holes;
4. a finite-corpus theorem for replay/compare gates and the proof obligation it leaves behind;
5. a sketch-level RosterSynth example and companion repository for runnable inspection.

The rest of the paper keeps that inspection path visible. Section 2 describes the setting that forced the problem and the harness we built. Section 3 names the failure. The lineage, artifact, role, lifecycle, and gate sections explain how a promoted case moves from SME correction into sketch, code or prompt, replay, and compare.

2 Originating setting

The enterprise deployment supplied the problem, not the finished method. Production failures could be loaded into a review surface; SMEs could correct concrete outputs; agents could repair code and prompts. Looking for a way to make those failures reusable led us back to CEGIS. The work here adapts that loop to agentic repair: load true failures from production, promote them into a golden set, and rerun the harness for each new scenario.

The deployed system joined eight pieces:

1. a custom Cursor extension for invoking repository actions from the IDE;
2. `cursor://`-style URL commands connecting review actions to local agent workflows;
3. an embedded web application where SMEs loaded production examples, inspected proposed remediations, and supplied corrections;
4. an execution runtime: AWS Lambda invoked through SQS, with model calls through the Bedrock Converse API;
5. a harness and test framework for fixtures, replay, semantic compare, golden corpus promotion, and provenance;
6. LLM-assisted conversion from SME corrections into input/output specifications;
7. a counterexample loop that promoted selected failures into a golden dataset;
8. non-developer operation of the loop for code and prompt implementations of a complex data-cleansing pipeline.

Lower formality was acceptable because the check stayed close to the people and artifacts that could verify it. SMEs inspected concrete outputs. The harness replayed each promoted scenario. The golden set made regressions visible. A correction was more than feedback on one row; it was a candidate specification change with a path into the repository.

The public artifact is a real but not real-world example for illustration. It keeps the shape of the work without exposing production data or client rules: a paper, fixtures, tests, provenance, and a small roster remediation case. The reusable loop is:

SME correction $\rightarrow$ input/output spec $\rightarrow$ promoted counterexample
 $\rightarrow$ sketch update $\rightarrow$ agent repair $\rightarrow$ replay/compare gate

The RosterSynth case gives readers a concrete artifact for inspecting how the method represents policy, encodes a promoted counterexample, and checks a repaired strategy.

3 Problem

Coding agents alter the economics of software change. They reduce the cost of producing a candidate patch. They also increase the number of candidate patches that look acceptable under incomplete specifications. That combination makes the specification boundary more important than the synthesis engine. Cheap patch generation helps only when the surrounding artifacts say which patch should be generated.

The sequence is recognizable:

1. A user gives an agent a domain task.
2. The prompt names the obvious state gap.
3. The agent finds a local patch that closes the visible gap.
4. A subject-matter expert recognizes a policy violation.
5. The violation was present in the domain, yet absent from the checkable artifacts.

The patch can be technically clean. The agent can follow local style. The tests can pass. The missing piece is the promoted counterexample: the artifact that says, “this plausible repair is forbidden for this reason, and future repairs must satisfy this rule.” Without that artifact, the team has learned something in conversation while leaving the repository nearly unchanged.

3.1 Tests need intent memory

A test is a predicate over behavior. A counterexample-supplemented sketch adds a memory of intent. The distinction matters when several outputs satisfy a visible predicate. A replay check can show that a proposed row closes a state delta. A semantic compare check can show that the row used the expected operation, target, date, or status. The sketch explains why those semantic fields matter.

The issue is the gap between what many tests encode and what a policy-heavy repair needs to preserve. Tests work well when the expected behavior is already explicit. A counterexample-supplemented sketch helps at the moment when the failure exposes a rule that had been implicit: it gives the rule a name, links it to the fixture that exposed it, and tells the next repair which tempting path to avoid.

Tests can also lose policy context when the failure is not linked to the rule it protects. A future maintainer or agent run may see only the failing assertion and patch around it. A sketch clause linked to the test records the policy ordering a later repair should preserve.

3.2 Prompts need repository anchoring

Run-local prompts can carry ambiguous intent in a way formal checks often cannot. They can state business rules, narrative criteria, exception handling, and domain language. That strength becomes fragile when the prompt lives in a chat transcript, notebook, model call, or agent session outside the repository’s review path.

The method still uses prompts. It treats them as one implementation surface. Oracle B is prompt-mediated completion constrained by the same sketch and corpus as deterministic code. Prompt text can encode narrative criteria; the resulting output remains subject to the same gate. The repository becomes the place where the prompt, the fixture, the sketch clause, and the check meet.

3.3 A repository-native version of the loop

Formal synthesis and CEGIS already cover a large design space. When the template, specification, and checker fit inside a formal language, a solver can search a constrained space and return a result with more semantic force than the repository gate used here. Counterexamples may come from a verifier, examples may guide inductive search, and oracles may steer component selection or specification refinement [16, 8].

The repository setting buys a different kind of evidence. Lower formality is acceptable when the rule can be checked by SMEs against a sketch and promoted counterexamples. Enterprise data-cleaning is one of several settings where this technique has promise: the common condition is SME-verifiable behavior before fully formalized rules. Some rules become deterministic code. Others stay in prompts while the team gathers enough counterexamples to encode them. Oracle B is not exempt from validation: prompt-mediated output is replayed and compared against the same promoted expectations.

This is a weaker result than formal synthesis. The trade is SME-verifiable evidence: sketch clauses, corrected examples, prompts, code anchors, checker fields, and generated provenance all point back to the rule the last tempting repair violated.

4 Lineage

4.1 Sketch and CEGIS

Solar-Lezama’s sketching work framed synthesis as collaboration between programmer insight and automated search: the programmer writes a partial program, leaving holes for a synthesizer to fill [14, 15]. Earlier sketching work already used counterexamples to refine candidates in finite-program synthesis settings [16]. The programmer supplies structure; the machine searches within it. CEGIS adds the validation loop: generate a candidate, validate it externally, and feed counterexamples back into the next synthesis step.

That loop carries an important discipline. A counterexample is more than a failing test case. It removes a region of the candidate space by showing a concrete input on which the current candidate violates the specification. In solver-backed settings, the synthesizer can use the growing counterexample set to search a formal space more intelligently than blind generate and check. Oracle-guided synthesis extends the picture by letting an oracle guide or validate choices in component-based synthesis [8].

4.2 Programming by example

Programming-by-example systems such as FlashMeta and PROSE synthesize DSL programs from examples using domain-specific deduction [12, 7]. Those systems show a different but related lesson: examples can be a specification surface when the target language and domain-specific operators are chosen carefully. The examples serve as constraints on a search procedure inside a language designed for that family of tasks.

4.3 Agentic coding and tests

Agentic coding work changes the repair surface again. Agent benchmarks such as SWE-bench evaluate whether language models can resolve real repository issues [9]. Tests-as-prompts work studies tests as both input and evaluation for LLM code generation [4]. These approaches already use repository artifacts and executable checks to steer agents. The method here is compatible with that direction. It adds a specific artifact boundary for policy-heavy repairs: every promoted counterexample should say which sketch clause changed, which tempting repair is rejected, which code or prompt surface handled the rule, and which gate check enforces it.

4.4 Prompt engineering, specs, and natural-language coding

Prompt programming and prompt-based learning treat natural-language prompts as a control surface for language models [13, 11]. Language-Oriented Programming makes the software claim: natural language can lower the barrier for non-experts who contribute to software projects [2]. Promptware engineering then asks teams to treat prompts like software artifacts with requirements, tests, debugging, evolution, and monitoring [3]. Copilot Workspace puts the same promise into product language: idea to code to software in natural language [5]. End-user software engineering is the older version of the same concern: domain experts often create or maintain computational artifacts without being professional software developers [10].

Spec-driven agent tools make a related move. Kiro uses requirements, design, and tasks to make model assumptions visible before implementation [17]. GitHub Spec Kit centers agentic development on a Spec–Plan–Tasks–Implement sequence and treats specs as structured context for agents [6]. That is useful language, but “spec” can still hide too much. Here, a spec means three things together: a sketch, promoted counterexamples, and a replay/compare harness. Specification by Example supplies the older lesson: concrete examples become living documentation only when they are executable enough to constrain future work [1].

Prompt text belongs in the paper because it is one mutable implementation surface in the loop. If a rule lives in a prompt today, the prompt needs the same pressure as code: a sketch clause, promoted examples, replay, semantic compare, and an owner. The sketch, corpus, anchors, and gate keep that surface reviewable.

4.5 Positioning boundary

Agentic synthesis against counterexample-supplemented sketches borrows from several traditions without treating any of them as a straw target. The sketch resembles a partial program. The corpus resembles the growing example set in CEGIS and programming by example. The gate resembles a bounded validator. The synthesizer is a coding agent with permission to edit source and prompts. The thing to inspect is the join among those artifacts.

The hole is shaped differently in an ordinary repository. In classical Sketch, the hole is usually inside a formal program and the solver searches a formal space. Here, the hole can be split across

Artifact	Reader question	Failure if blurred
Sketch	What rule should future repairs preserve?	The prompt drifts and the rule stays hidden.
Corpus	Which cases have specification authority?	Examples become anecdotes, not promoted constraints.
Anchors	What local code shape must the agent preserve?	The agent invents a second convention.
Implementation	Which code or prompt handles the hole?	A paper rule cannot be traced to the surface that applies it.
Gate	What behavior is executable now?	The claim floats above the repository instead of being checked.

Table 1: Artifact responsibilities from a reader’s point of view.

code, fixtures, prompts, review comments, and generated evidence. A rule may need one deterministic branch, one prompt instruction, one fixture, and one semantic compare field. Counterexamples keep those pieces tied to the rule they are supposed to preserve.

	Sketch / CEGIS	Programming by example	Repository loop
Human input	Partial program with holes	Examples inside a DSL	Sketch, SME corrections, counterexamples, anchors
Synthesizer	Solver-backed search	DSL learner	Coding agent editing code and prompts
Failure signal	Counterexample from validator	Missing or inconsistent example	Gate failure or SME correction
Validation	Formal or bounded decision procedure	Example consistency	Replay plus semantic compare over $\mathcal{E}$
Result boundary	Solver-relative result	DSL/example consistency	Finite-corpus predicate satisfaction under repo semantics

Read the table as a positioning map, not a taxonomy of synthesis or agentic coding. Each column names one family resemblance: partial programs, example-driven constraints, repository repair, and executable validation. The contribution is the repository join: sketch clause, promoted counterexample, code or prompt surface, and gate check in one inspectable path.

5 Artifacts

The method assigns each artifact one job. The sketch records strategy. The corpus records promoted cases. Known-code anchors record local implementation shape. The gate records the current executable semantics. Keeping those jobs separate makes failures easier to interpret: when a gate fails, the maintainer can ask whether the sketch, corpus, anchor, implementation, or checker needs to change.

Definition 1 (Sketch). *A sketch $\mathcal{S}$ is a human-editable artifact that fixes the permitted strategy space and names holes. It may contain prose, tables, pseudo-code, type shapes, policy order, abstention rules, and examples of forbidden repairs. It persists in the repository and changes when a counterexample reveals missing policy.*

A sketch should answer these questions for future maintainers and agent runs:

1. Which operations exist?

2. Which operation has priority when several repairs close the same state gap?
3. Which fields are policy-bearing?
4. Which decisions belong in deterministic code?
5. Which decisions belong in prompt-mediated narrative completion?
6. Which cases force abstention or escalation?

The sketch is deliberately less formal than a full specification. Its job is to hold the policy shape before every detail can be encoded. That incompleteness is useful when the team knows a rule exists but has not yet decided whether it belongs in code, prompt text, fixture data, or human review.

Definition 2 (Corpus). *The corpus $\mathcal{E} = \{(x_i, y_i)\}_{i=1}^n$ is the finite set of promoted examples the current gate must satisfy. Each input x_i is a concrete scenario or fixture. Each expected output y_i is the semantic repair the gate enforces.*

The corpus is authoritative because promotion is deliberate. Raw examples can be noisy. SME corrections can contain ambiguity. A promoted counterexample should meet an explicit maintainer-reviewed threshold: it names a plausible wrong output, states the violated rule, and is encoded in the gate. Promotion turns a case from observed behavior into specification pressure.

Definition 3 (Known-code anchor). *A known-code anchor $\mathcal{K}$ is an existing source artifact that constrains the agent’s implementation style: an API, result type, parser convention, error shape, fixture format, test helper, reference implementation, or prompt structure. Anchors reduce the risk of parallel local inventions.*

Anchors matter because agentic repair can accidentally create a second convention beside the first. A sketch names the policy. Anchors name the shape the codebase already uses to carry that policy. They are especially important when an agent can satisfy a local test with a new result shape that no downstream caller expects.

Definition 4 (Gate). *A gate $\mathcal{G}$ is the executable validation bundle that evaluates a candidate strategy $\mathcal{H}$ over every promoted case in $\mathcal{E}$. In this method the gate has two layers: replay and semantic compare.*

Replay checks whether the proposed output changes the world state according to the encoded state predicate. Semantic compare checks whether the output used the encoded policy-bearing fields from the promoted expectation. The two layers separate state repair from policy repair. That separation matters because the most dangerous repair can be the one that makes the arithmetic look right while choosing the wrong operation.

6 Roles

The method assigns distinct jobs to humans, agents, and checks. The separation is practical. SMEs know which output is acceptable. Maintainers decide which corrections become specification. Agents repair implementation surfaces. Gates report whether the current strategy satisfies the encoded corpus.

6.1 Subject-matter expert

The SME supplies corrections on concrete examples through a review interface. Their review action says: this proposed output is wrong; the correct output is this; the reason is this domain rule. In the originating harness, the system helped convert that correction into an input/output spec. In the general method, the workflow needs some conversion path from SME judgment to repository artifact.

The SME’s correction has three parts:

1. the observed input;
2. the corrected output;
3. the explanation that distinguishes the corrected output from the tempting wrong one.

The explanation matters because it seeds the sketch update and directs repair toward the general rule the fixture exemplifies. A corrected row alone can invite memorization. A corrected row plus a reason tells the next agent why the row matters.

6.2 Developer or platform maintainer

The developer supplies the rails: runtime, review UI, extraction path, fixtures, replay, semantic compare, tests, and provenance. In the originating harness, that meant Lambda invoked through SQS, Bedrock Converse for model calls, repository-local commands, and gates that could be run again after promotion.

The developer does not hand-author every domain rule. SMEs and the harness control the rule flow: corrections enter through the review interface, the harness turns them into candidate specification artifacts, and promoted cases become part of $\mathcal{E}$. Developer ownership is over the machinery and anchors: whether the runtime runs, whether the gate checks the right fields, whether the output shape stays compatible, and whether a promoted rule remains inspectable.

6.3 Coding agent

The agent repairs the artifact under constraints. It may edit deterministic code, prompt text, fixtures, or tests. Its job is synthesis inside the boundary set by $\mathcal{S}$, $\mathcal{E}$, $\mathcal{K}$, and $\mathcal{G}$.

The agent should receive three forms of context before editing:

1. the sketch clause involved;
2. the counterexample and tempting wrong patch;
3. the known-code anchors that define output shape and local conventions.

That context changes the repair problem. The agent is no longer asked only to make a failing signal green. It is asked to preserve a specific rule, reject a known tempting patch, and keep the codebase’s existing shape.

6.4 Oracle A: deterministic code

Oracle A handles encodable policy. It should be deterministic, reviewable source code: grouping, filtering, ordering, type construction, deterministic abstention, and domain calculations. The agent can maintain it because the rule has been made concrete enough to encode.

6.5 Oracle B: prompt-mediated completion

Oracle B handles narrative or under-encoded policy: ambiguous notes, human descriptions, client-specific language, or judgment that belongs in a model prompt for the current stage of the system. Oracle B remains constrained by the sketch and checked by the same replay/compare gate. Its output is judged against the SME-approved corpus instead of treated as authority.

6.6 Hybrid resolver

A hybrid resolver tries Oracle A first, then routes abstentions to Oracle B. This keeps deterministic policy on the reliable path and leaves narrative completion explicit. The split also gives the team a migration path: a rule can begin in prompt-mediated completion, then move into deterministic code once enough counterexamples make it precise.

7 Counterexample lifecycle

The counterexample lifecycle is the core process. It turns an observed failure into a repository artifact the next agent must respect. Each stage has a separate owner so the loop keeps SME judgment, maintainer promotion, agent repair, and executable validation distinct.

7.1 Observation

A concrete case enters the review surface. It can come from production data, a synthetic fixture, a regression report, or a reviewer. The system proposes an output. The SME accepts or corrects it. At this stage the case is evidence, not yet specification.

7.2 Correction

A correction names the desired output and the rule behind it. When possible, it also names the tempting wrong repair. That is where tacit domain knowledge becomes actionable: an expert says which nearby output would have passed replay, and why policy still rejects it.

7.3 Spec extraction

The harness converts the correction into an input/output spec. An LLM can help draft this spec, especially when the SME explanation is prose. The maintainer decides whether that draft is precise enough to promote; the gate later enforces the encoded expectation. This step is a translation step, not an automatic source of truth.

7.4 Promotion

Promotion makes the case part of $\mathcal{E}$. Promotion is a specification change. The promoted case must contain enough information to reject the tempting wrong repair. A case that only records the correct answer can still leave the rule ambiguous.

Promotion also creates maintenance work. The sketch may need a new clause, the comparer may need a new field, the fixture may need a clearer name, and the failure message may need to tell the next agent where to look. If those pieces are missing, the corpus grows without making the repair problem clearer.

7.5 Sketch revision

The sketch changes with the corpus. The revised sketch should state the general rule that the counterexample exposed and the fixture answer that demonstrates it. This protects the next agent pass from memorizing one row. The sketch is the place where the team names the rule; the fixture is the place where the team records the concrete pressure that forced the rule into view.

7.6 Repair

The agent repairs Oracle A, Oracle B, fixtures, or tests under the revised sketch. Known-code anchors constrain result shape and local convention. The repair should close the promoted case while preserving earlier promoted cases. Otherwise the loop has only moved the failure.

7.7 Gate

The gate runs replay and semantic compare over all promoted cases. A passing gate establishes the Section 9 result for the current strategy, corpus, checkers, fixtures, and golden rows. If any of those change, the gate must run again.

8 Gate semantics

The gate is the executable semantics of the current corpus. It has two layers because state repair and semantic repair fail differently. Keeping those layers separate gives the next agent a better error signal: a replay failure says the output failed the encoded state repair; a compare failure says the output repaired state while violating an encoded policy field.

Definition 5 (Replay). *Replay applies a candidate output r to an input state x and checks whether the proposed change closes the domain gap. Replay is domain-specific: it may compute balances, apply updates, recalculate derived fields, or simulate downstream effects.*

Definition 6 (Semantic compare). *Semantic compare checks fields that replay leaves open: operation kind, selected entity, work date, issue type, status mutation, route, priority, or any other policy-bearing field encoded by the comparer. A row is compare-valid when it matches the promoted expectation on those fields.*

Definition 7 (Gate pass). *For a strategy $\mathcal{H}$ and corpus $\mathcal{E}$, the gate passes iff every $(x_i, y_i) \in \mathcal{E}$ yields $r_i = \mathcal{H}(x_i)$ such that replay accepts (x_i, r_i) and semantic compare accepts (y_i, r_i) .*

Replay is responsible for encoded state-repair failures. Compare is responsible for encoded policy-field violations. A useful gate usually needs both because the two checks reject different wrong repairs. In the roster example, an append repair can close the hours math while choosing the wrong operation; replay accepts the state change and compare rejects the policy violation.

8.1 Design rules for gates

A gate should be deterministic, local, and specific. It should run locally, report the exact promoted case that failed, and preserve the difference between replay failure and compare failure. The failure message should tell the next agent which rule to inspect.

Good gates have these properties:

1. **Promoted corpus coverage.** Every promoted case runs.

Candidate behavior	Replay	Compare	Interpretation
Open state gap	reject	not reached	Candidate fails encoded state repair.
Closed state, wrong op	accept	reject	Candidate repairs state but violates a policy field.
Closed state, right op	accept	accept	Case satisfies $\mathcal{E}$ under current repo semantics.

Table 2: Replay and semantic compare answer different reader questions.

2. **Failure specificity.** Replay and compare failures are reported separately.
3. **Deterministic CI path.** Gate-critical checks do not depend on a live model call.
4. **Source links.** Each failure can be traced to sketch clause, scenario, and check.
5. **Tempting-patch rejection.** The gate includes at least one check that fails the previously plausible wrong repair.

These rules are design constraints rather than formal assurances. A gate that omits a field cannot enforce that field. A comparer with a bug can certify the wrong behavior. Gate design is specification work, and checker quality belongs in the limitations rather than behind the word correctness.

9 Finite-corpus correctness

The proof obligation is intentionally bounded. A passing gate establishes $\mathcal{E}$ -correctness relative to the current replay function, semantic comparer, fixtures, and golden rows. When a new failure is promoted into $\mathcal{E}$, it creates a larger proof obligation; after the gate passes again, the same bounded result applies to the expanded corpus.

Definition 8 ($\mathcal{E}$ -correctness). *A strategy $\mathcal{H}$ is $\mathcal{E}$ -correct with respect to replay function R and compare predicate C when, for every $(x_i, y_i) \in \mathcal{E}$, $R(x_i, \mathcal{H}(x_i)) = \text{true}$ and $C(y_i, \mathcal{H}(x_i)) = \text{true}$.*

Lemma 1 (Replay soundness relative to the implementation). *If replay returns true for (x, r) , then r satisfies the state-repair condition encoded by replay for x .*

Proof. Replay is the executable definition of state repair for the gate. It evaluates the candidate output against the input state and returns true exactly when the encoded state-repair predicate holds. A true replay result entails the state-repair condition the gate implements. $\square$

Lemma 2 (Compare soundness relative to the corpus). *If semantic compare returns true for expected output y and candidate output r , then r agrees with y on every policy-bearing field encoded by the comparer.*

Proof. The comparer enumerates the checked policy-bearing fields and fails on mismatch. A true compare result means every encoded policy-bearing field matched the promoted expectation. $\square$

Theorem 1 (Finite-corpus gate soundness). *If $\mathcal{G}$ passes for strategy $\mathcal{H}$ on corpus $\mathcal{E}$, then $\mathcal{H}$ is $\mathcal{E}$ -correct with respect to the replay and compare semantics encoded by the repository.*

Repair	Visible state effect	Policy meaning	Gate result
Append adjustment	hours delta closed	gap compensation	replay accepts; compare rejects
Cancel duplicate	hours delta closed	duplicate removed	replay accepts; compare accepts

Table 3: The worked example separates state repair from policy repair in one small case.

Proof. By definition, $\mathcal{G}$ iterates over every $(x_i, y_i) \in \mathcal{E}$, computes $r_i = \mathcal{H}(x_i)$, and requires replay and compare to pass. By replay soundness, each accepted r_i satisfies the state-repair condition encoded by replay for x_i . By compare soundness, each accepted r_i agrees with y_i on every encoded policy-bearing field. Thus every promoted corpus case satisfies both predicates, which is exactly $\mathcal{E}$ -correctness. $\square$

Scope. The theorem covers the promoted corpus, the encoded replay checker, the encoded semantic comparer, and the current golden rows. It says nothing about unpromoted failures, unencoded policy fields, future model behavior, or private deployment behavior outside the public artifact. Checker bugs require checker fixes. Wrong golden rows require human correction. The result is narrow enough to audit and strong enough to guide agentic repair: a maintainer can see which predicate passed, which corpus it covered, and which boundary remains open.

10 Worked example at sketch level

RosterSynth is the paper’s sketch-level worked example. It illustrates the loop through one roster remediation case. The companion repository contains the implementation details for running the example, including source, fixtures, tests, and generated provenance. The paper stays at sketch level so the reader can see the method before tracing file paths.

10.1 Scenario

A subject-matter expert sees a roster gap and two active bookings that explain it. The visible arithmetic suggests an adjustment. The policy points to duplicate cancellation. The important feature of the example is the ambiguity: two repairs can close the same apparent state gap, and only one follows the rule.

The tempting local interpretation is simple: if badge hours and scheduled hours disagree, append an adjustment that closes the delta. That patch can make the numbers balance. The policy interpretation is different: when duplicate active bookings with the same relevant shape explain the gap, cancel the selected duplicate. The difference is small in output shape and large in meaning.

10.2 Tempting repair

The tempting repair adds an adjustment that closes the hours balance. Replay accepts the state change because the arithmetic gap disappears. Semantic compare rejects the repair because the selected operation is wrong. The method retains this failure because the patch succeeds under one visible predicate while violating the policy-bearing field.

10.3 Sketch rule

The counterexample promotes a sketch rule: duplicate active bookings with the same relevant shape on the pay-window end date cancel the selected duplicate when that cancellation closes the gap. The sketch now tells the next agent which operation has priority over the adjustment. It also tells the maintainer what future changes must preserve: the operation order, the selected duplicate, and the fields that distinguish cancellation from compensation.

10.4 Gate lesson

The example illustrates the paper’s central gate lesson in one roster remediation case. Replay asks, “did the proposed output repair the encoded state gap?” Compare asks, “did it use the encoded rule the sketch requires?” The method needs both questions because a state-valid repair can still be a policy violation.

10.5 Artifact boundary

The paper stops at sketch level. The companion repository supplies source, fixtures, tests, and provenance for concrete-path inspection. This split lets the paper focus on the technique while the repository carries implementation evidence. A reader who wants the exact command path can follow the README; a reader who wants the claim boundary can stay with the theorem and the sketch-level example.

11 Applying the method

A lightweight version of the workflow can precede a full enterprise harness. The practical objective is to make every agent repair answer three questions: which rule is being used, which example exposed it, and which gate would reject the tempting wrong repair. Teams can start with ordinary repository files and add automation only where the loop already works.

11.1 Start with the sketch

Write the intended strategy before asking the agent for code. The sketch should name the available operations, the priority order, the known holes, and the abstention cases. Even an incomplete sketch can be useful if it gives later failures a place to attach. The first sketch is a working surface, not a claim that the policy is complete.

11.2 Collect examples before polishing code

Examples are useful when they distinguish competing repairs. A happy-path example shows what success looks like. A counterexample shows which plausible success should be rejected. The second kind is especially useful for agentic repair because it teaches the boundary between “green” and “right under this rule.”

11.3 Name known-code anchors

Point the agent at existing result types, parser behavior, fixture format, error shape, and prompt conventions. This reduces the chance that the agent creates a second local convention just to satisfy the current case. Anchors turn style and shape from tacit repository knowledge into explicit repair context.

11.4 Separate replay from compare

Build replay and compare as separate checks. A single assertion that says “the output looks right” hides the distinction between state repair and policy repair. Separate failures teach the next agent more: replay failure means the state predicate stayed broken; compare failure means the state predicate passed while a policy-bearing field diverged from the promoted expectation.

11.5 Promote failures deliberately

Every failure should be triaged. Some failures expose fixture bugs. Some expose checker bugs. Some expose missing policy. The last category belongs in $\mathcal{E}$ and $\mathcal{S}$. Promotion is the moment where a case becomes part of the specification.

Deliberate promotion asks for more than a corrected answer. It asks what tempting patch the case rejects, which sketch clause should change, which field the comparer must check, and which failure message will guide a future repair. That extra work is the price of turning a single failure into reusable specification pressure.

12 Companion artifact

Open the companion repository before trusting the argument too much. It contains the sanitized RosterSynth slice, generated provenance, and tests that exercise the sketch-level story. The README points to the PDF, the source, the evidence appendix, the sketch, the promoted corpus, Oracle A code, hybrid resolver, Oracle B prompt path, and tests.

The artifact has three jobs:

1. let developers run the example and watch the gate pass or fail;
2. let academics trace the finite-corpus claim into code;
3. show future users which files would change in another domain.

That boundary keeps the paper from pretending to be the repo. The paper explains the method. The README routes readers. Example docs show the concrete flow. Generated appendices expose provenance for readers who want to audit the path from sketch to check.

13 Discussion

13.1 Tests name behavior; sketches name intent

Tests tell the current run what passed. A sketch clause tells the next repair why the failure mattered. When a failing case joins $\mathcal{E}$, the sketch records the rule future agent passes must preserve. That is the useful temporal distinction: behavior now, intent next time.

13.2 Prompts implement one surface of the contract

Prompts handle narrative holes, ambiguous notes, and policy text that has not yet become code. They are not outside the method. The sketch and corpus constrain them, and the gate checks their outputs. In this loop, prompting is one implementation surface beside source code, fixtures, and tests.

13.3 SME corrections become specification inputs

In the originating harness, SMEs corrected outputs through the review interface instead of opening tickets for every specification change. The harness turned those corrections into candidate specs and counterexamples. Maintainers still decided what entered $\mathcal{E}$. The SME contribution was the pressure: “this output is wrong, and here is the domain rule it missed.”

13.4 Promotion quality is a scaling constraint

A large pile of examples is weak guidance when the rule stays ambiguous. A smaller corpus of sharp counterexamples can be more useful because each case rejects a specific tempting patch. The promotion question is concrete: what wrong repair would this case still allow if we left it out?

14 Limitations

1. **Finite corpus.** The theorem covers $\mathcal{E}$. New failures must be promoted and the gate must pass again before the bounded result covers them. The method does not claim universal correctness across unseen cases.
2. **Checker quality.** Replay and compare inherit the limits of their encoded semantics. A buggy checker can certify the wrong behavior until a maintainer repairs the checker and replays the promoted corpus.
3. **Golden quality.** A wrong golden row makes the gate enforce the wrong behavior. Promotion needs review by someone who understands the domain rule and can distinguish a corrected answer from a corrected specification.
4. **Promotion quality.** A promoted case should distinguish a rule from a recorded answer. Weak promotion gives future agents examples to memorize; strong promotion gives them rules to preserve.
5. **Sketch quality.** Hidden prose rules still require human review. The method gives those rules a repository home, then relies on maintainers to keep that home current as new counterexamples arrive.
6. **Prompt drift.** Live model behavior can drift. The method treats live output as something to review and promote, not as a stable proof artifact.
7. **Enterprise evidence boundary.** The public companion demonstrates the control structure on sanitized fixtures. The proprietary harness motivates the method and remains outside the public artifact; the public evidence supports the loop and states the enterprise deployment boundary.

The limitations are not cleanup after the method. They mark where the proof stops and where a person still owns the result: promotion, checker design, golden review, prompt review, and private deployment claims.

15 Conclusion

The useful part is not that a coding agent can write a patch. The useful part is the trail it has to leave before anyone trusts the patch: the sketch clause, the promoted case, the code or prompt surface, the replay check, and the semantic compare check.

Agentic synthesis against counterexample-supplemented sketches treats the agent as a synthesizer inside that constraint environment. A human writes the sketch and promotes failures. The SME corrects concrete outputs. The agent edits code and prompts. The gate states the bounded result for the current corpus.

That loop came from an enterprise setting where SMEs had the domain knowledge and agents had the editing surface. The reusable technique is the repository join between them. Each promoted counterexample adds an executable check against a known wrong repair, and each passing gate says only what it can honestly say: the current strategy satisfies the encoded replay and semantic-comparison predicates for the current promoted corpus.

References

- [1] Gojko Adzic. *Specification by Example: How Successful Teams Deliver the Right Software*. Manning Publications, 2011.
- [2] Amin Beheshti. Natural language-oriented programming (NLOP): Towards democratizing software creation. In *2024 IEEE International Conference on Software Services Engineering, SSE 2024*, pages 258–267. IEEE, 2024.
- [3] Zhenpeng Chen, Chong Wang, Weisong Sun, Xuanzhe Liu, Jie M. Zhang, and Yang Liu. Promptware engineering: Software engineering for prompt-enabled systems, 2025.
- [4] Yi Cui. Tests as prompt: A test-driven-development benchmark for llm code generation, 2025.
- [5] Thomas Dohmke. Github copilot workspace: Welcome to the copilot-native developer environment. GitHub Blog, 2024. Published April 29, 2024.
- [6] GitHub. What is spec-driven development? GitHub Spec Kit Documentation, 2026. Accessed July 1, 2026.
- [7] Sumit Gulwani, Oleksandr Polozov, and Rishabh Singh. Program synthesis. *Foundations and Trends in Programming Languages*, 4(1–2):1–119, 2017.
- [8] Susmit Jha, Sumit Gulwani, Sanjit A. Seshia, and Ashish Tiwari. Oracle-guided component-based program synthesis. In *Proceedings of the 32nd ACM/IEEE International Conference on Software Engineering, ICSE 2010*, pages 215–224, 2010.
- [9] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R. Narasimhan. SWE-bench: Can language models resolve real-world github issues? In *The Twelfth International Conference on Learning Representations*, 2024.
- [10] Andrew J. Ko, Robin Abraham, Laura Beckwith, Alan Blackwell, Margaret Burnett, Martin Erwig, Chris Scaffidi, Joseph Lawrance, Henry Lieberman, Brad Myers, Mary Beth Rosson, Gregg Rothmel, Mary Shaw, and Susan Wiedenbeck. The state of the art in end-user software engineering. *ACM Computing Surveys*, 43(3), 2011.
- [11] Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *ACM Computing Surveys*, 55(9):1–35, 2023.

- [12] Oleksandr Polozov and Sumit Gulwani. Flashmeta: A framework for inductive program synthesis. In *Proceedings of the 2015 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications*, OOPSLA 2015, pages 107–126. ACM, 2015.
- [13] Laria Reynolds and Kyle McDonell. Prompt programming for large language models: Beyond the few-shot paradigm. *arXiv preprint arXiv:2102.07350*, 2021.
- [14] Armando Solar Lezama. *Program Synthesis by Sketching*. PhD thesis, EECS Department, University of California, Berkeley, December 2008.
- [15] Armando Solar-Lezama. Program sketching. *International Journal on Software Tools for Technology Transfer*, 15(5–6):475–495, 2013.
- [16] Armando Solar-Lezama, Liviu Tancau, Rastislav Bodik, Sanjit A. Seshia, and Vijay A. Saraswat. Combinatorial sketching for finite programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS 2006, pages 404–415. ACM Press, 2006.
- [17] Nikhil Swaminathan and Deepak Singh. Introducing kiro. Kiro Blog, 2025. Published July 14, 2025.